

# Temporal Difference Learning

# Temporal Difference Learning

- Break up the episode and learn directly from actual experience from the environment.
- TD is model free: no knowledge of transitions/rewards.
- Learn from incomplete episodes, by bootstrapping (i.e. update our guess of the value function from a certain step on).

# MC and TD

- Goal: Given  $\pi$ , learn  $q_\pi$  online from experience.
- Incremental every-visit Monte Carlo
  - Update estimate of the value  $Q(S_t, A_t)$  towards **actual** return  $G_t$

$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(G_t - Q(S_t, A_t))$$

- TD(0):
  - Update estimate of the value  $V(S_t)$  towards the **estimated** return  $Q(S_t, A_t)$  by:


$$Q(S_t, A_t) \leftarrow Q(S_t, A_t) + \alpha(R_{t+1} + \gamma Q(S_{t+1}, A_{t+1}) - Q(S_t, A_t))$$

# MC vs TD

- TD can learn before the final episode, unlike MC.
- TD works in non-terminating environments, unlike MC.
- TD exploits the Markov property (hence more effective in Markov environments).
- MC does not exploit the Markov property (more effective in non-Markov environments).


# Monte-Carlo backup

$$V(S_t) \leftarrow V(S_t) + \alpha (G_t - V(S_t))$$




# Temporal Difference backup

$$V(S_t) \leftarrow V(S_t) + \alpha (R_{t+1} + \gamma V(S_{t+1}) - V(S_t))$$



# Dynamic programming backup

$$V(S_t) \leftarrow \mathbb{E}_{\pi} [R_{t+1} + \gamma V(S_{t+1})]$$



# On and Off-Policy Learning

- On-policy learning:
  - "Learn on the job"
  - Learn about  $\pi$  by sampling experience from  $\pi$ .
- Off-policy learning:
  - "Shadow" someone.
  - Learn about policy  $\pi$  from experience sampled from policy  $\pi'$ .
  - Learn about *optimal* policy while following an *exploratory* policy.



# On-Policy Learning: Sarsa

Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
Repeat (for each episode):  
    Initialize  $S$   
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
    Repeat (for each step of episode):  
        Take action  $A$ , observe  $R, S'$   
        Choose  $A'$  from  $S'$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
         $Q(S, A) \leftarrow Q(S, A) + \alpha[R + \gamma Q(S', A') - Q(S, A)]$   
         $S \leftarrow S'; A \leftarrow A';$   
    until  $S$  is terminal

# Example: Frozen Lake

- [https://gym.openai.com/evaluations/eval\\_HmYHTSY5QYOkI77yacPLkg](https://gym.openai.com/evaluations/eval_HmYHTSY5QYOkI77yacPLkg)


# Off-Policy Learning: Q-Learning

```
Initialize  $Q(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A}(s)$ , arbitrarily, and  $Q(\text{terminal-state}, \cdot) = 0$   
Repeat (for each episode):  
  Initialize  $S$   
  Repeat (for each step of episode):  
    Choose  $A$  from  $S$  using policy derived from  $Q$  (e.g.,  $\epsilon$ -greedy)  
    Take action  $A$ , observe  $R, S'$   
     $Q(S, A) \leftarrow Q(S, A) + \alpha [R + \gamma \max_a Q(S', a) - Q(S, A)]$   
     $S \leftarrow S'$ ;  
  until  $S$  is terminal
```

## Example: Frozen Lake

- [https://gym.openai.com/evaluations/eval\\_j8yzfQ4BQ9O6dI5Pr7SrtQ](https://gym.openai.com/evaluations/eval_j8yzfQ4BQ9O6dI5Pr7SrtQ)

# Q-Learning vs Sarsa



# Cliffworld - Results



# Afterstates

